

# Contents

|          |   |          |
|----------|---|----------|
| <b>1</b> | <b>Introduction</b>                                 | <b>3</b> |
| 1.1      | Common Directories Structure . . . . .              | 3        |
| 1.2      | Symbolic Links . . . . .                            | 5        |
| 1.3      | Ways to find installed software . . . . .           | 5        |
| 1.4      | Scripting and the shell . . . . .                   | 6        |
| 1.4.1    | Files and Directories . . . . .                     | 6        |
| 1.4.2    | Manipulating Files . . . . .                        | 7        |
| 1.4.3    | Pipes and Redirection . . . . .                     | 7        |
| 1.4.4    | Variables and Quoting . . . . .                     | 8        |
| 1.4.5    | Common Filter Commands . . . . .                    | 9        |
| 1.4.6    | Regular Expressions . . . . .                       | 11       |
| 1.4.7    | Other Commands . . . . .                            | 12       |
| 1.5      | Controlling Processes . . . . .                     | 13       |
| 1.5.1    | PID - Process ID Number . . . . .                   | 13       |
| 1.5.2    | PPID: Process Parent ID . . . . .                   | 13       |
| 1.5.3    | UID and EUID: Real and Effective user ID . . . . .  | 13       |
| 1.5.4    | GID and EGID: Real and Effective Group ID . . . . . | 13       |
| 1.5.5    | Signals . . . . .                                   | 13       |
| 1.5.6    | KILL: Send Signals . . . . .                        | 14       |
| 1.5.7    | Process States . . . . .                            | 14       |
| 1.5.8    | PS: Monitor Processes . . . . .                     | 14       |
| 1.5.9    | Dynamic Monitoring . . . . .                        | 15       |
| 1.5.10   | The /Proc Filesystem . . . . .                      | 15       |
| 1.5.11   | Trace Signals and System Calls . . . . .            | 15       |



# Chapter 1

## Introduction

### 1.1 Common Directories Structure

Common utilization of some directories in a Linux system:

- `/` : The root directory where the file system begins. In most cases the root directory only contains subdirectories.
- `/boot` : This is where the Linux kernel and boot loader files are kept. The kernel is a file called *vmlinuz*.
- `/etc` : The `/etc` directory contains the configuration files for the system. All of the files in `/etc` should be text files. Points of interest:
  - `/etc/passwd` : The `passwd` file contains the essential information for each user. It is here that users are defined.
  - `/etc/fstab` : The `fstab` file contains a table of devices that get mounted when your system boots. This file defines your disk drives.
  - `/etc/hosts` : This file lists the network host names and IP addresses that are intrinsically known to the system.
  - `/etc/init.d` : This directory contains the scripts that start various system services typically at boot time.
- `/bin`, `/usr/bin` : These two directories contain most of the programs for the system. The `/bin` directory has the essential programs that the system requires to operate, while `/usr/bin` contains applications for the system's users.
- `/sbin` `/usr/sbin` : The `sbin` directories contain programs for system administration, mostly for use by the superuser.
- `/usr` : The `/usr` directory contains a variety of things that support user applications. Some highlights:

- */usr/share/X11* : Support files for the X Window system.
- */usr/share/dict*: Dictionaries for the spelling checker.
- */usr/share/doc* : Various documentation files in a variety of formats.
- */usr/share/man* : The man pages are kept here.
- */usr/src* : Source code files. If you installed the kernel source code package, you will find the entire Linux kernel source code here.
- */usr/local* : */usr/local* and its subdirectories are used for the installation of software and other files for use on the local machine. What this really means is that software that is not part of the official distribution (which usually goes in */usr/bin*) goes here.

When you find interesting programs to install on your system, they should be installed in one of the */usr/local* directories. Most often, the directory of choice is */usr/local/bin*.

- */var* : The */var* directory contains files that change as the system is running. This includes:
  - */var/log* : Directory that contains log files. These are updated as the system runs.
  - */var/spool* : This directory is used to hold files that are queued for some process, such as mail messages and print jobs.
- */lib* : The shared libraries (similar to DLLs in that other operating system) are kept here.
- */home* : is where users keep their personal work. In general, this is the only place users are allowed to write files.
- */root* : This is the superuser's home directory.
- */tmp* : is a directory in which programs can write their temporary files.
- */dev* : The */dev* directory is a special directory, since it does not really contain files in the usual sense. Rather, it contains devices that are available to the system. In Linux (like Unix), devices are treated like files. You can read and write devices as though they were files. For example */dev/fd0* is the first floppy disk drive, */dev/sda* is the first hard drive. All the devices that the kernel understands are represented here.
- */proc* : The */proc* directory is also special. This directory does not contain files. In fact, this directory does not really exist at all. It is entirely virtual. The */proc* directory contains little peep holes into the kernel itself. There are a group of numbered entries in this directory that correspond to all the processes running on the system. In addition, there are a number of named entries that permit access to the current configuration of the system. Many of these entries can be viewed. Try viewing */proc/cpuinfo*. This entry will tell you what the kernel thinks of your CPU.

- */media*, */mnt* : The */media* directory is used for mount points. The process of attaching a device to the file system tree is called mounting. For a device to be available, it must first be mounted. When your system boots, it reads a list of mounting instructions in the file */etc/fstab*, which describes which device is mounted at which mount point in the directory tree. This takes care of the hard drives, but you may also have devices that are considered temporary, such as CD-ROMs, thumb drives, and floppy disks. Since these are removable, they do not stay mounted all the time. The */media* directory is used by the automatic device mounting mechanisms found in modern desktop oriented Linux distributions. To see what devices and mount points are used use the ***mount*** command.

## 1.2 Symbolic Links

Symbolic links are a special type of file that points to another file. With symbolic links, it is possible for a single file to have multiple names. Here's how it works: Whenever the system is given a file name that is a symbolic link, it transparently maps it to the file it is pointing to.

To create symbolic links, use the ***ln*** command.

## 1.3 Ways to find installed software

**Whatis** command is helpful to get brief information about Linux commands or functions. **Whatis** command displays man page single line description for command that matches string passed as a command line argument to **what**is command. Example:

```
luca@scipppane:~$ whatis echo
echo (1) - display a line of text
```

If we want to search Linux commands or functions information using wild card, then **what**is command gives **-w** option. Example:

```
luca@scipppane:~$ whatis -w "ab*"
abort (3) - cause abnormal process termination
abs (3) - compute the absolute value of an integer
```

If we want to search Linux commands or functions information using regular expressions, then **what**is command gives **-r** option. Example (all commands ending with ab):

```
luca@scipppane:~$ whatis -r "ab$"
anacrontab (5) - configuration file for anacron
baobab (1) - A graphical tool to analyze disk usage
crontab (1) - maintain crontab files for individual users (Vixie Cron)
crontab (5) - tables for driving cron
dvgrab (1) - Capture DV or MPEG-2 Transport Stream (HDV) video and audio da
```

|   |   |
|---|---|
| <code>fstab</code> (5)                    | – static information about the filesystems    |
| <code>inittab</code> (5)                  | – init daemon configuration                   |
| <code>pmxab</code> (1)                    | – a MusiXTeX preprocessor                     |
| <code>scilab</code> (1)                   | – manual page for Unknown argument: – version |
| <code>swab</code> (3)                     | – swap adjacent bytes                         |
| <code>tc-stab</code> (8)                  | – Generic size table manipulations            |
| <code>XChangeActivePointerGrab</code> (3) | – grab the pointer                            |
| <code>XcmsCIELab</code> (3)               | – Xcms color structure                        |
| <code>XtAddGrab</code> (3)                | – redirect user input to a modal widget       |
| <code>XtRemoveGrab</code> (3)             | – redirect user input to a modal widget       |

**which** command to find out if a relevant binary is already in your search path.

```
luca@scipppane:~$ which gcc
/usr/bin/gcc
```

If **which** can't find the command you're looking for, try **whereis**; it searches a broader range of system directories and is independent of your shell's search path. Example:

```
luca@scipppane:~$ whereis open
open: /bin/open /usr/share/man/man2/open.2.gz /usr/share/man/man1/open.1.gz
```

If we want to locate binary of Linux command, use **"-b"** option.

```
luca@scipppane:~$ whereis -b open
open: /bin/open
```

Another alternative is the incredibly useful **locate** command, which consults a precompiled index of the filesystem to locate filenames that match a particular pattern.

## 1.4 Scripting and the shell

### 1.4.1 Files and Directories

To find the name of the working directory, use the ***pwd*** command.

To list the files in the working directory, use the ***ls*** command.

Common options:

- ***ls -l*** : list the files in the working directory in long format
- ***ls -la*** : list all files (even ones with names beginning with a period character, which are normally hidden) in long format

To change your working directory, use the ***cd*** command. To do this, type ***cd*** followed by the pathname of the desired working directory.

The "." notation refers to the working directory itself and the ".." notation refers to the working directory's parent directory.

Common options:

- *cd* : followed by nothing will change the working directory to your home directory
- *cd ~ username* : will change the working directory to the home directory of the specified user

### 1.4.2 Manipulating Files

Key commands are:

- *cp* : copy files and directories
- *mv* : move or rename files and directories
- *rm* : remove files and directories
- *mkdir* : create directories

Wildcards allow to select filenames based on patterns of characters.

You can use wildcards with any command that accepts filename arguments.

| Wildcard               | Meaning  |
|------------------------|--|
| *                      | Matches any character  |
| ?                      | Matches any single character   |
| [ <i>characters</i> ]  | Matches any character that is a member of the set <i>characters</i> . The set of characters may also be expressed as a POSIX character class such as one of the following:<br><br>:alnum: Alphanumeric characters<br>:alpha: Alphabetic characters<br>:digit: Numerals<br>:upper: Uppercase alphabetic characters<br>:lower: Lowercase alphabetic characters |
| [! <i>characters</i> ] | Matches any character that is not a member of the set <i>characters</i>  |

### 1.4.3 Pipes and Redirection

Every process has at least three communication channels available to it: "standard input" (STDIN), "standard output" (STDOUT), and "standard error" (STDERR).

The kernel sets up these channels on the process's behalf, so the process itself

doesn't necessarily know where they lead. They might connect to a terminal window, a file, a network connection, or a channel belonging to another process, to name a few possibilities. Most commands accept their input from STDIN and write their output to STDOUT. They write error messages to STDERR. This convention lets you string commands together like building blocks to create composite pipelines.

The shell interprets the symbols `<`, `>`, and `»` as instructions to reroute a command's input or output to or from a file. A `<` symbol connects the command's STDIN to the contents of an existing file. The `>` and `»` symbols redirect STDOUT; `>` replaces the file's existing contents, and `»` appends to them. For example, the command

```
echo "This is a test message" > /tmp/mymessage
```

stores a single line in the file `/tmp/mymessage`, creating the file if necessary. To redirect both STDOUT and STDERR to the same place, use the `>&` symbol. To redirect STDERR only, use `2>`.

To connect the STDOUT of one command to the STDIN of another, use the `|` symbol, commonly known as a pipe. Some examples:

```
ps -ef | grep httpd
```

To execute a second command only if its precursor completes successfully, you can separate the commands with an `&&` symbol. For example,

```
$ lpr /tmp/t2 && rm /tmp/t2
```

removes `/tmp/t2` if and only if it is successfully queued for printing.

Conversely, the `||` symbol executes the following command only if the preceding command fails.

In a script, you can use a backslash to break a command onto multiple lines, helping to distinguish the error-handling code from the rest of the command pipeline:

```
cp --preserve --recursive /etc/* /spare/backup \
|| echo "Did NOT make backup"
```

For the converse effect, multiple commands combined onto one line, you can use a semicolon as a statement separator.

#### 1.4.4 Variables and Quoting

Variable names are unmarked in assignments but prefixed with a dollar sign when their values are referenced. For example:

```
luca@scippane:~$ etcdirc='/etc'
luca@scippane:~$ echo $etcdirc
/etc
```

Do not put spaces around the `=` symbol or the shell will mistake your variable name for a command name.



When referencing a variable, you can surround its name with curly braces to clarify to the parser and to human readers where the variable name stops and other text begins; for example, `${etcdir}` instead of just `$etcdir`.

Variable names are case sensitive.

Environment variables are automatically imported into bash's variable namespace, so they can be set and read with the standard syntax. Use **export** *var-name* to promote a shell variable to an environment variable. Commands for environment variables that you want to set up at login time should be included in your `/.profile` or `/.bash_` profile file.

The shell treats strings enclosed in single and double quotes similarly, except that double-quoted strings are subject to globbing (the expansion of filename matching metacharacters such as `*` and `?`) and variable expansion. For example:

```
luca@scippane:~$ myland="Pennsylvania Duch"
luca@scippane:~$ echo "I speak $myland"
I speak Pennsylvania Duch
luca@scippane:~$ echo 'I speak $myland'
I speak $myland
```

Back quotes, also known as back-ticks, are treated similarly to double quotes, but they have the additional effect of executing the contents of the string as a shell command and replacing the string with the command's output. For example,

```
luca@scippane:~$ echo "There are `wc -l /etc/passwd` lines in the password file"
There are 88 /etc/passwd lines in the password file
```

### 1.4.5 Common Filter Commands

Any well-behaved command that reads STDIN and writes STDOUT can be used as a filter (that is, a component of a pipeline) to process data.

**cut:** *separate lines into fields*

The cut command prints selected portions of its input lines.

```
cut -d: -f7 < /etc/passwd | sort -u
```

The cut command picks out the path to each user's shell from `/etc/passwd`. The list of shells is then sent through `sort -u` to produce a sorted list of unique values.

**sort:** *sort lines*

sort sorts its input lines.

| Opt | Meaning  |
|-----|--|
| -b  | Ignore leading whitespace                        |
| -f  | Case insensitive sorting                         |
| -k  | Specify the columns which form the sort key      |
| -n  | Compare fields as integer numbers                |
| -r  | Reverse sort order                               |
| -t  | Set field separator (The default is white space) |
| -u  | Output unique records only                       |

Examples:

```
sort -t: -k3,3 -n /etc/group
root:x:0:
bin:x:1:daemon
daemon:x:2:
```

```
sort -t: -k3,3 /etc/group
root:x:0:
bin:x:1:daemon
users:x:100:
```

**uniq:** *print unique lines*

`uniq` is similar in spirit to `sort -u`, but it has some useful options that `sort` does not emulate: **-c** to count the number of instances of each line, **-d** to show only duplicated lines, and **-u** to show only nonduplicated lines. The input must be presorted, usually by being run through `sort`.

For example the command below shows that 50 users have `/bin/bash` as their login shell and that 18 have `/bin/false` and so on.

```
cut -d: -f7 /etc/passwd | sort | uniq -c
    50 /bin/bash
    18 /bin/false
     2 /bin/sh
     1 /bin/sync
    17 /usr/sbin/nologin
```

**wc:** *count lines, words, and characters*

Counting the number of lines, words, and characters in a file is another common operation, and the **wc** (word count) command is a convenient way of doing this. Run without options, it displays all three counts:

```
luca@scippane:~$ wc /etc/passwd
 88  114 5032 /etc/passwd
```

In the context of scripting, it is more common to supply a **-l**, **-w**, or **-c** option to make `wc`'s output consist of a single number.

**tee:** *copy input to two places* A command pipeline is typically linear, but it's often helpful to tap into the data stream and send a copy to a file or to the terminal window. You can do this with the **tee** command, which sends its stan-

ward input both to standard out and to a file that you specify on the command line. For example:

```
find /lib/modules -name core | tee /dev/tty | wc -l
```

prints both the pathnames of files named core and a count of the number of core files that were found.

**head** and **tail**: *read the beginning or end of a file*

These commands display ten lines by default, but you can include a command-line option to specify how many lines you want to see.

**tail** also has a nifty **-f** option that's particularly useful for sysadmins. Instead of exiting immediately after printing the requested number of lines, **tail -f** waits for new lines to be added to the end of the file and prints them as they appear; great for monitoring log files.

**grep**: *search text*

**grep** searches its input text and prints the lines that match a given pattern. Its name is based on the `g/regular-expression/p` command from the old **ed** editor that came with the earliest versions of UNIX (and still does).

Like most filters, **grep** has many options, including **-c** to print a count of matching lines, **-i** to ignore case when matching, and **-v** to print nonmatching (rather than matching) lines. Another useful option is **-l** (lowercase L), which makes **grep** print only the names of matching files rather than printing each line that matches.

For example:

```
$ sudo grep -l mdadm /var/log/*
/var/log/auth.log
/var/log/syslog.0
```

shows that log entries from mdadm have appeared in two different log files.

### 1.4.6 Regular Expressions

Regular expressions are supported by most modern languages. They are also used by UNIX commands such as **grep** and **vi**.

The filename matching and expansion performed by the shell when it interprets command lines such as **wc -l \*.pl** is not a form of regex matching. It's a different system called "shell globbing" and it uses a different and simpler syntax. Regular expressions reached the apex of their power and perfection in **Perl**.

**The Matching Process** Code that evaluates a regular expression attempts to match a single given text **string** to a single given **pattern**.

For the matcher to declare success, the entire search pattern must match a contiguous section of the search text. However, the pattern can match at any position.

**Literal Characters** In general, characters in a regular expression match themselves.

**Special Characters** The table below shows the meanings of some common special symbols that can appear in regular expressions. These are just the basics; there are many, many more.

| Symbol                      | What it matches or does  |
|-----------------------------|--|
| *                           | Matches any character  |
| [ <i>chars</i> ]            | Matches any character from a given set                         |
| [^ <i>chars</i> ]           | Matches any character NOT in a given set                       |
| ^                           | Matches the beginning of a line                                |
| &                           | Matches the beginning of a line                                |
| \w                          | Matches any word character                                     |
| \s                          | Matches any whitespace character                               |
| \d                          | Matches any digit  |
|                             | Matches either the element to its left or the one to its right |
| ( <i>expr</i> )             | Limits scope, groups elements, allows matches to be captured   |
| ?                           | Allows zero or one match of the preceding element              |
| *                           | Allows zero, one, or many matches of the preceding element     |
| +                           | Allows one or more matches of the preceding element            |
| { <i>n</i> }                | Matches exactly <i>n</i> instances of the preceding element    |
| { <i>min</i> ,}             | Matches at least <i>min</i> instances (note the comma)         |
| { <i>min</i> , <i>max</i> } | Matches any number of instances from <i>min</i> to <i>max</i>  |

Many special constructs, such as + and |, affect the matching of the "thing" to their left or right. In general, a "thing" is a single character, a subpattern enclosed in parentheses, or a character class enclosed in square brackets.

If you want to limit the scope of the vertical bar, enclose the bar and both things in their own set of parentheses. For example,

```
I am the (walrus|egg man)\.
```

matches either "I am the walrus." or "I am the egg man.". This example also demonstrates escaping of special characters (here, the dot).

### 1.4.7 Other Commands

**ls:** *List files and directories*

With option *-l* lists the file in long format

With option *-la* lists all files

With option *-d* list directory entries instead of contents, and do not dereference symbolic links

**file:** *Identify file type*

**ln:** *Create symbolic link*

**cp:** *Copy files and directories*

With option *-R* copy the contents of *dir1* into *dir 2*

**mv:** *Move or rename files and directories*

**rm:** *Remove files and directories*

With option *-r* removes an entire directory

**mkdir:** *Create directories*

## 1.5 Controlling Processes

A process consists of an address space and a set of data structures within the kernel. The address space is a set of memory pages that the kernel has marked for the process's use.

### 1.5.1 PID - Process ID Number

The kernel assigns a unique ID number to every process. Most commands and system calls that manipulate processes require you to specify a PID to identify the target of the operation. PIDs are assigned in order as processes are created.

### 1.5.2 PPID: Process Parent ID

When a process is cloned, the original process is referred to as the parent, and the copy is called the child. The PPID attribute of a process is the PID of the parent from which it was cloned.

### 1.5.3 UID and EUID: Real and Effective user ID

A process's UID is the user identification number of the person who created it, or more accurately, it is a copy of the UID value of the parent process. Usually, only the creator (aka the "owner") and the superuser can manipulate a process. The EUID is the "effective" user ID, an extra UID used to determine what resources and files a process has permission to access at any given moment.

### 1.5.4 GID and EGID: Real and Effective Group ID

The GID is the group identification number of a process. The EGID is related to the GID in the same way that the EUID is related to the UID in that it can be "upgraded" by the execution of a `setgid` program.

### 1.5.5 Signals

Signals are process-level interrupt requests.

When a signal is received, one of two things can happen. If the receiving process has designated a handler routine for that particular signal, the handler is called with information about the context in which the signal was delivered. Otherwise, the kernel takes some default action on behalf of the process. Some signals with which all administrators should be familiar are:

| No | Name | Description          | Default   |
|----|------|----------------------|-----------|
| 1  | HUP  | Hangup               | Terminate |
| 2  | INT  | Interrupt            | Terminate |
| 3  | QUIT | Quit                 | Terminate |
| 9  | KILL | Kill                 | Terminate |
| 11 | SEGV | Segmentation Fault   | Terminate |
| 15 | TERM | Software Termination | Terminate |

### 1.5.6 KILL: Send Signals

As its name implies, the **kill** command is most often used to terminate a process. **kill** can send any signal, but by default it sends a TERM. **kill** can be used by normal users on their own processes or by root on any process. The syntax is:

```
kill [-signal] pid
```

where *signal* is the number or symbolic name of the signal to be sent (as shown in Table above) and *pid* is the process identification number of the target process. The command

```
kill -9 pid
```

"guarantees" that the process will die because signal 9, KILL, cannot be caught.

Under Linux, **killall** kills processes by name.

### 1.5.7 Process States

You need to be aware of the four execution states listed in the table below:

| State    | Meaning   |
|----------|---|
| Runnable | The process can be executed                       |
| Sleeping | The process is waiting for some resources         |
| Zombie   | The process is trying to die                      |
| Stopped  | The process is suspended (not allowed to execute) |

A runnable process is ready to execute whenever CPU time is available.

It has acquired all the resources it needs and is just waiting for CPU time to process its data.

Sleeping processes are waiting for a specific event to occur.

Zombies are processes that have finished execution but have not yet had their status collected.

Stopped processes are administratively forbidden to run.

### 1.5.8 PS: Monitor Processes

**ps** is the system administrator's main tool for monitoring processes.

**ps** can show the PID, UID, priority, and control terminal of processes. It also gives information about how much memory a process is using, how much CPU

time it has consumed, and its current status.

On Linux and AIX, you can obtain a useful overview of all the processes running on the system with **ps aux**. The **a** option means to show all processes, **x** means to show even processes that don't have a control terminal, and **u** selects the "user oriented" output format.

Another useful set of arguments for Linux and AIX is **lax**, which provides more technical information (**l** selects the "long" output format).

### 1.5.9 Dynamic Monitoring

**top** is a free utility that runs on many systems and provides a regularly updated summary of active processes and their use of resources.

By default, the display updates every 10 seconds. The most CPU-consumptive processes appear at the top.

### 1.5.10 The /Proc Filesystem

The Linux versions of **ps** and **top** read their process status information from the **/proc** directory, a pseudo-filesystem in which the kernel exposes a variety of interesting information about the system's state.

### 1.5.11 Trace Signals and System Calls

Linux lets you directly observe a process with the **strace** command, which shows every system call the process makes and every signal it receives.